How much each model thought at each setting

Thinking tokens per answer on the eleven photos. Bar: average. Line: how much it varies from one answer to the next.

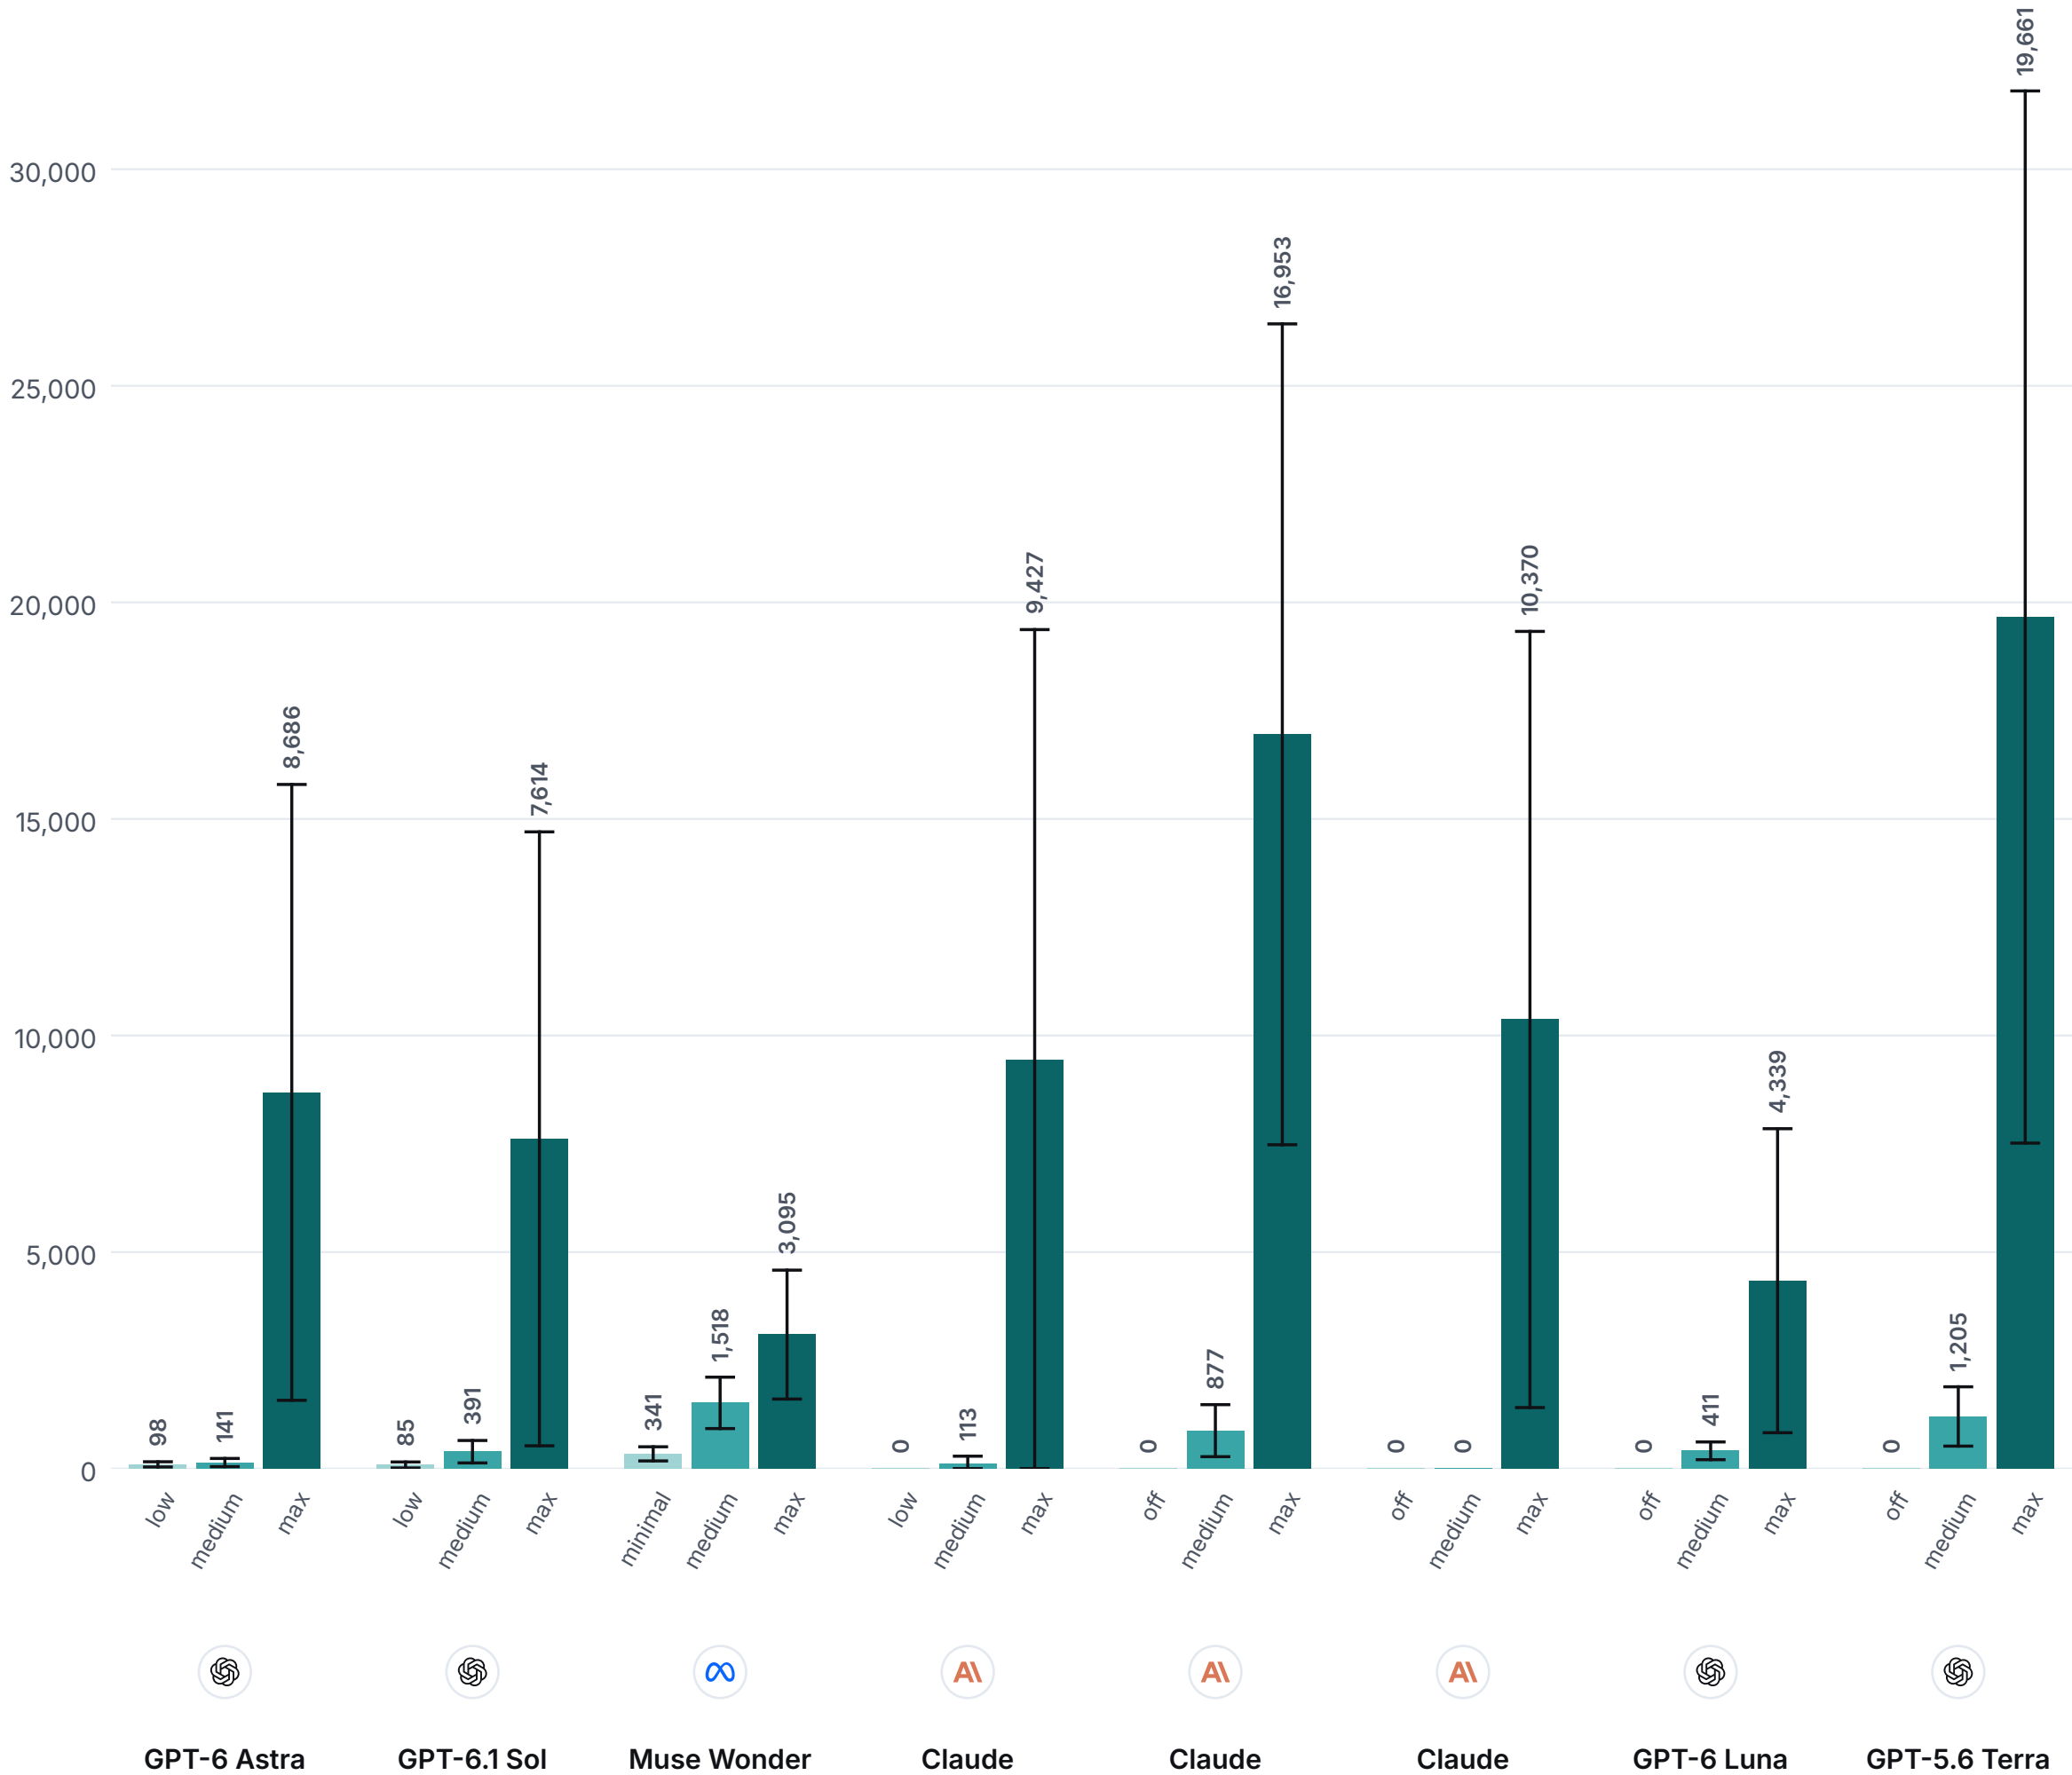

33 answers per bar (11 photos, 3 tries; Terra at max 32). Line: one standard deviation either side, cut at 0; none when every answer is the same.

Counts as each provider reports them. A Claude answer may use at most 32,000 tokens, thinking included; the ChatGPT plan set no such limit.