How much each model thinks before pointing

Thinking tokens per answer on the eleven photos. Bar: average. Line: how much it varies from one answer to the next.

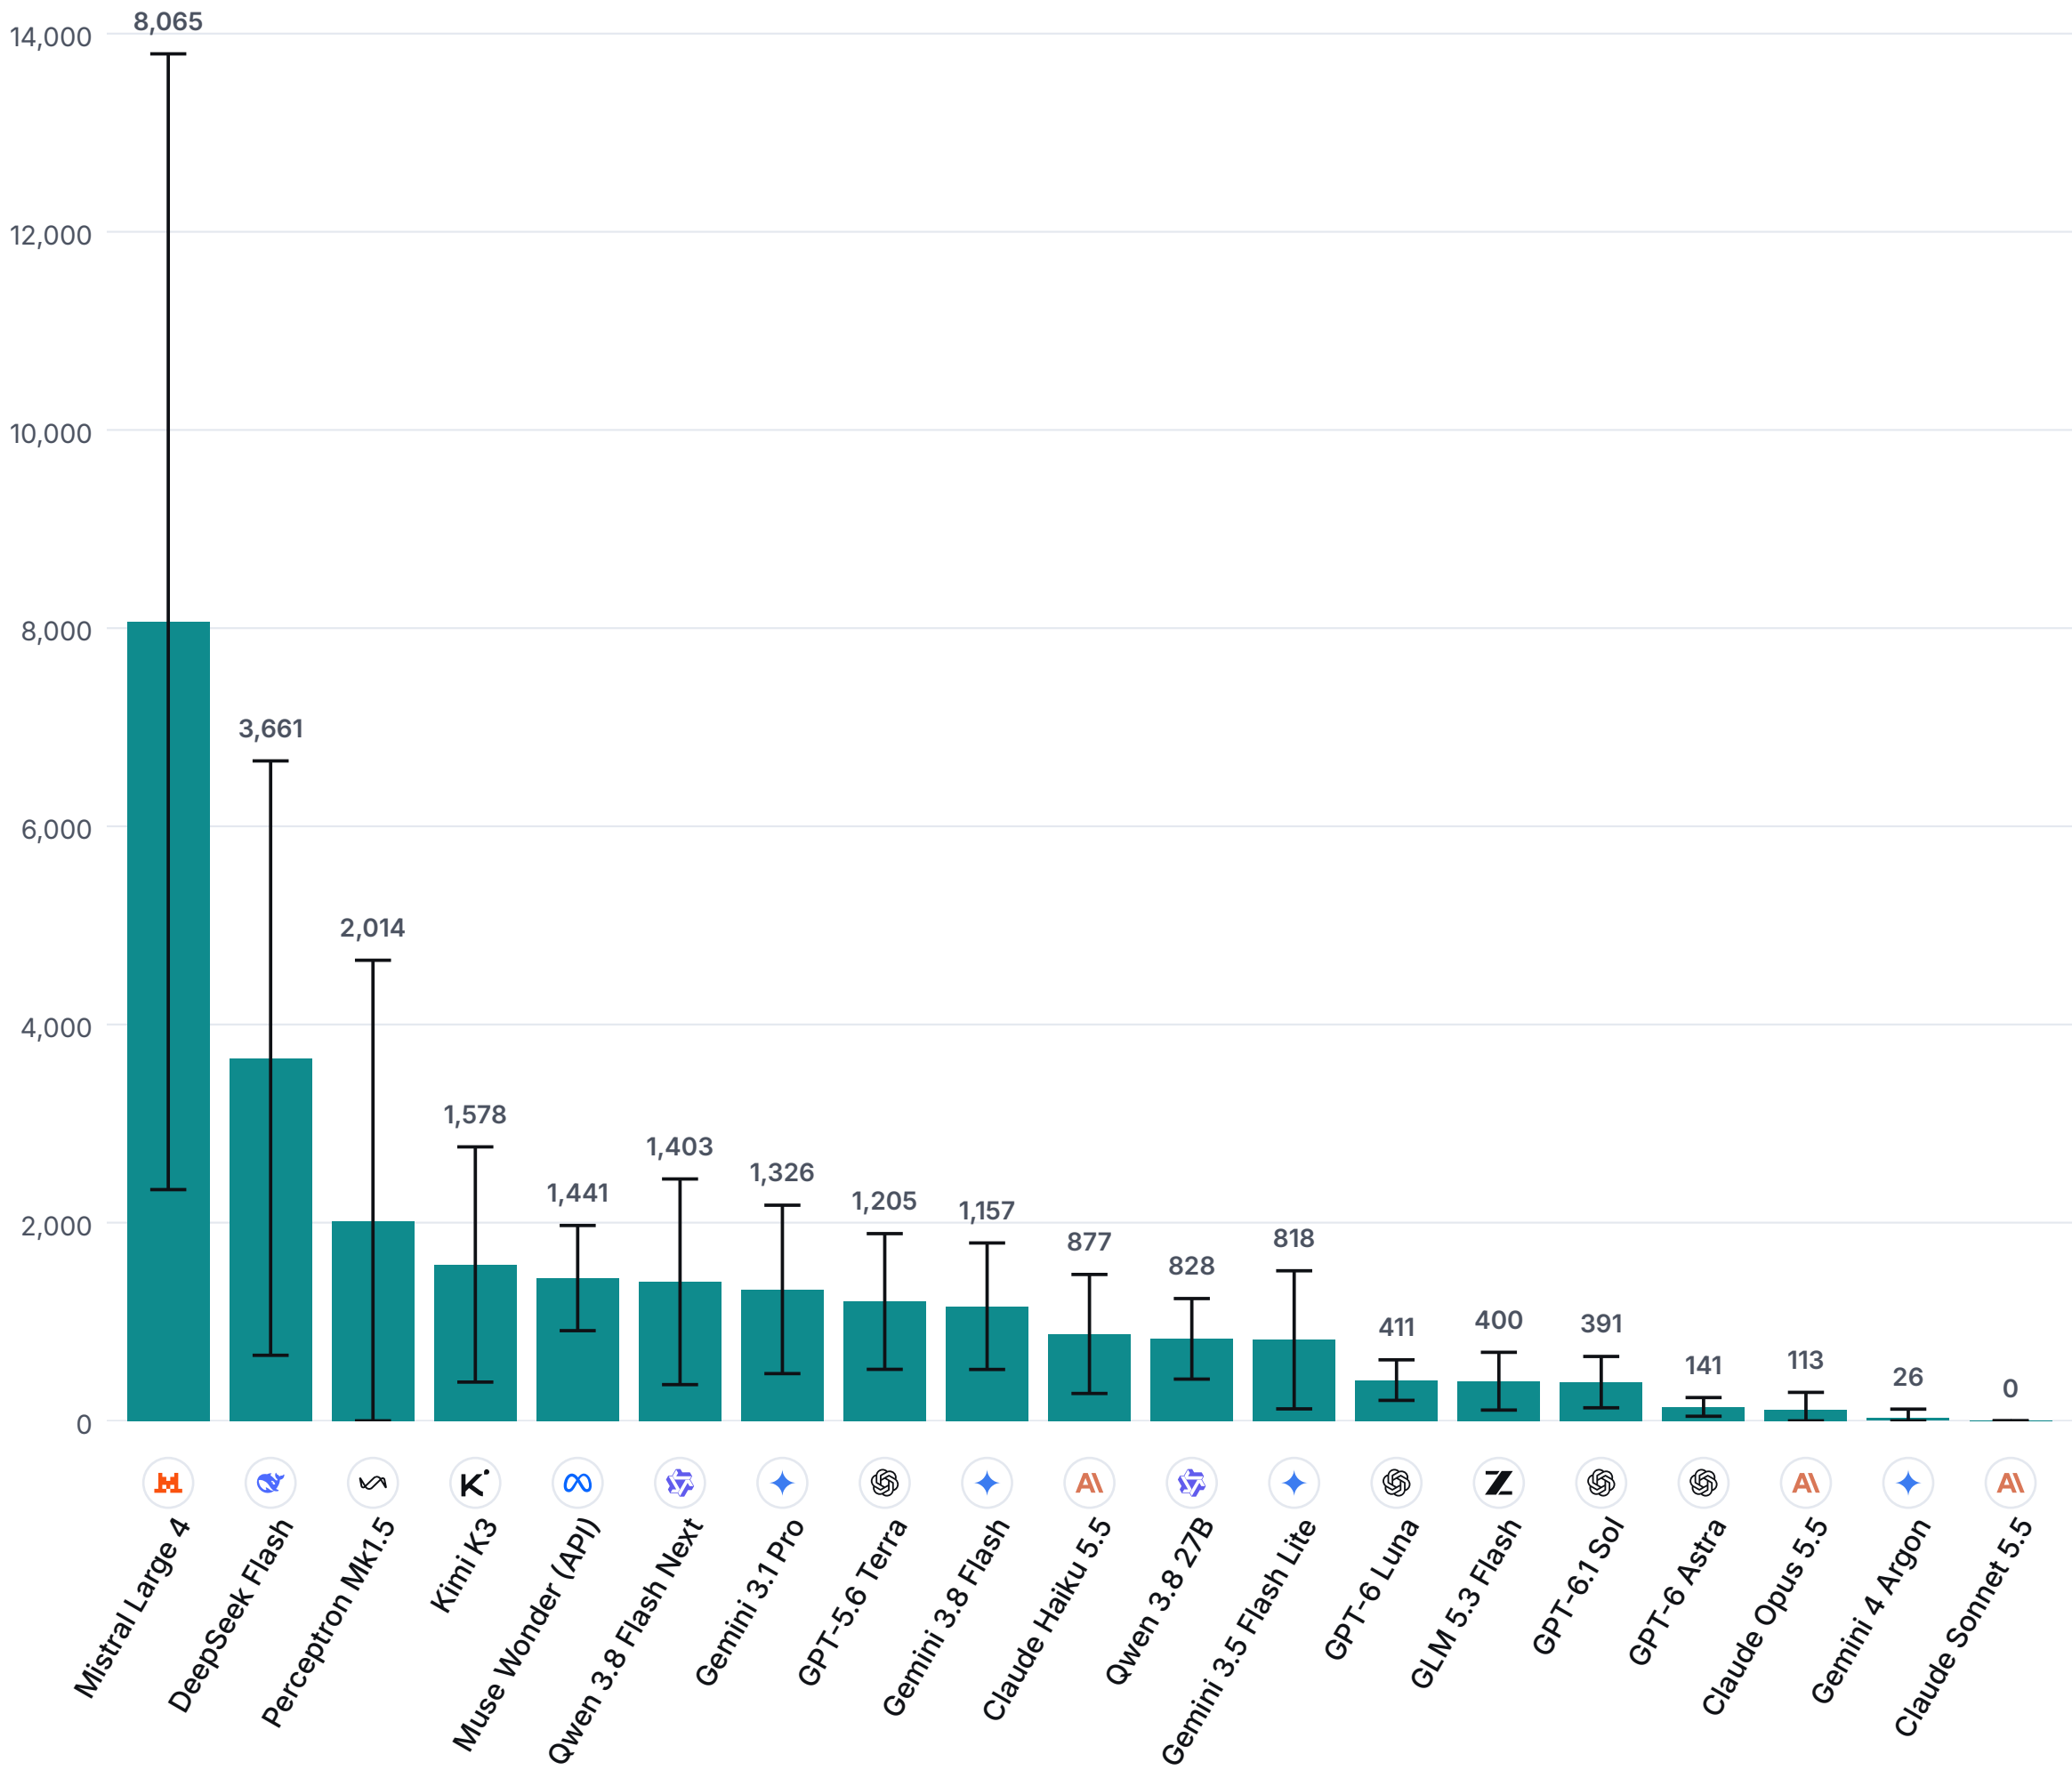

627 answers: 19 language models, 11 photos, 3 tries, at the medium thinking setting used on the photo chart (GLM, Kimi and Mistral: their high). Line: one standard deviation either side, cut at 0. Counts as each provider reports them. Local models were run without thinking.